

CSU34031 Advanced Telecommunications

Project-1 - Senán d'Art - 17329580

Introduction

This project is to:

- Develop a web proxy that can handle HTTP, HTTPS and Websocket connections.
- Support caching of HTTP packets
- Support the dynamic blocking of URLs through a management console
- Support multithreaded operation

I chose to use NodeJS and JavaScript to accomplish these goals.

High-Level Overview

For every connection to the proxy server (multiple occurrences possible for each user):

- If the requested **URL is not blacklisted**:
 - If the connection **uses TLS or is a Websocket**:
 - Pipe all packets from server to client
 - Pipe all packets from client to server
 - Notes:
 - There is no reason to parse this data further
 - Else the connection **does not use TLS and is not a websocket**:
 - If the **packet is cached already**:
 - If the **cached item is not expired**:
 - Update the age of the packet and send it as response to the client
 - Else the **cached item is expired**:
 - Remove from cache and continue as if it was never cached
 - Else **packet is not cached**:
 - Send request to server
 - On server response:
 - If the **response packet is cache-able** (based on header params):
 - If **response is chunked**:
 - Send each chunk to the user as it arrives
 - Temporarily store each chunk
 - When transmission is complete -> add all chunked data to cache
 - Else **response is not chunked**
 - Send response packet to client
 - Add packet to cache
 - Else the **response packet is not cache-able**:
 - Pipe the response to the user
 - Else the requested **URL is blacklisted**:
 - Send a **403 FORBIDDEN** response to the client
 - Close the connection

Installation

Requires:

- NodeJS 12.13.1+
- npm (if for some reason it wasn't bundled with Node)

Set up:

1. Open a terminal in the directory containing `proxy.js` and `package.json`
2. Run `npm install` (may require sudo)
3. Run `npm start`

Usage

Commands:

- `block <domain>` - adds specified domain to blocklist
- `unblock <domain>` - removes domain from blocklist
- `cache` - enables caching
- `nocache` - disables caching
- `verbose` - prints all connections to console
- `noverbose` - disables printing of all connections to console
- `timing` - print timing data of cache hits/misses
- `notiming` - disables printing of cache hit/miss timing data
- `showsaving` - shows how many bytes have been served from cache
- `showcachesize` - shows the current cache size

Multithreading

NodeJS operates based on events and callbacks. When an event occurs and has an event handler, that event is handled by a thread. In the case of the `net` library I am using, it will choose a thread from the process threadpool. To multithread the program I assigned the process up to 1000 threads in its threadpool. This means that the proxy can run with up to 1000 threads. Since each thread runs asynchronously this would be enough for many connections with low latency.

Testing

In order to test the performance of the proxy and the effect caching has on responsiveness, I loaded some websites with and without caching.

All tests were conducted with Firefox 73.0.1 and repeated 3 times with a restart of the proxy used between each to clear the cache. The `noverbose` flag was enabled for all tests to reduce clutter.

I ran the same test on the following sites:

- `www.example.com`
 - No-Cache
 - Webpage
 - `99ms, 100ms, 99ms`
 - Average: `~100ms`
 - Favicon

- 532ms, 550ms, 565ms
 - Average: ~545ms
 - Cache
 - Webpage
 - 0ms, 1ms, 1ms
 - Average: ~0.5ms
 - Favicon
 - 550ms, 541ms, 511ms
 - Average: ~540ms
 - This one had a header requesting not to be cached
 - Bandwith saved: 3,162 Bytes
- www.writephononline.com
 - Notes:
 - There are lots of requests for various assets on this page so the following are averages
 - Most of the requests under the hood are actually HTTPS or not cache-able
 - No-Cache
 - 1880ms, 2200ms, 2280ms
 - Average: ~2100ms
 - Cache
 - 1520ms, 1280ms, 1400ms
 - Average: ~1400ms
 - Bandwitdh saved: 3,818 Bytes

Code

```
process.env.UV_THREADPOOL_SIZE = 1000 //maximum number of threads
require('events').EventEmitter.defaultMaxListeners = 20 //maximum
listeners per event
const net = require('net')
const fs = require('fs')
const path = require('path')
const { exec } = require('child_process')
const server = net.createServer()
const port = 4000
const stdin = process.openStdin()
const blockListName = 'blockList.json'
const blockListPath = path.join(__dirname, blockListName)
let verbose = true
let caching = true
let timing = true
let currentCacheSize = 0
let bytesSavedFromNetwork = 0 //how much bandwidth has been saved
let suppressErrs = true

/**
 * Read and return blocklist
 * @return {{blockedURLs:Array}}
 */
let readBlockList = () => {
```

```
    return JSON.parse(fs.readFileSync(blockListPath, 'utf8'))
  }
  let blockList = readBlockList()

  let cache = []

  //No clients connected
  server.on('close', () => {
    console.log(`All clients disconnected`)
  })

  //Something broke
  server.on('error', (err) => {
    console.error({ ERROR: err })
    exec('npm start', (err, stdout, stderr) => {
      if (err) {
        console.error(err)
      } else {
        console.log(`stdOUT: ${stdout}`)
        console.log(`stdERR: ${stderr}`)
      }
    })
  })

  //new connection to server
  server.on('connection', (clientProxyConn) => {
    //create the connection

    clientProxyConn.once('data', (data) => {
      let theData = data.toString()
      let reqData = getAddrAndPort(theData)
      //if it's blacklisted, send 403 and kill connection
      if (blockList.blockedURLs.includes(reqData.host)) {
        clientProxyConn.write('HTTP/1.1 403 FORBIDDEN\r\n\r\n')
        clientProxyConn.end()
        clientProxyConn.destroy()
        if (verbose)
          console.log({
            Message: 'Connection Blocked',
            Hostname: reqData.host,
            Port: reqData.port,
            HTTPS: reqData.isHTTPS
          })
        return
      }

      let toServerConn = net.createConnection({
        host: reqData.host,
        port: reqData.port
      }, () => {
        //if is HTTPS, confirm connection
        //else send the request to the server
        if (reqData.isHTTPS || !caching) {
          if (reqData.isHTTPS)
```

```

        clientProxyConn.write('HTTP/1.1 200 OK\r\n\r\n')
        //Don't manually handle subsequent data streams, this is
easier, faster and uses less memory
        //readableSrc.pipe(writableDest)
        clientProxyConn.pipe(toServerConn).pipe(clientProxyConn)
    } else {
        if (isWebsocketRequest(data)) { //don't cache websockets,
or headers that request 'no-cache'
            console.log(`Not caching websocket request for:
${reqData.rawURL}`)

clientProxyConn.pipe(toServerConn).pipe(clientProxyConn)
        } else { //need to manually handle chunked data
            let cachedRes = getFromCache(reqData.rawURL)
            let startTime = Date.now()
            if (!cachedRes) { //response not already cached
                toServerConn.write(data)
                let dataWhole = []
                let isChunked = false
                let cacheableResponse = false
                let checkedCachability = false
                toServerConn.on('data', (resData) => {
                    if (!checkedCachability) { //need to check if
we can cache it
                        cacheableResponse =
isCacheableResponse(resData)
                        checkedCachability = true
                    }

                    if (!cacheableResponse) { //not cacheable
                        console.log(`Packet from
'${reqData.rawURL}', does not allow caching`)
                        clientProxyConn.write(resData)

clientProxyConn.pipe(toServerConn).pipe(clientProxyConn)
                    } else {
                        if (isChunked ||
resData.toString().includes('Transfer-Encoding: chunked\r\n')) { //for
chunked data need to cache differently
                            console.log('Chunked data incoming!')
                            clientProxyConn.write(resData)

                            dataWhole.push(resData)
                            if (!isChunked)
                                isChunked = true

                            if (resData.toString().slice(-5) ==
'0\r\n\r\n') { //end of chunked encoding
                                addToCache(dataWhole[0],
reqData.rawURL, dataWhole.splice(0, 1))
                            }
                        } else {
                            // console.log('multiple requests in
same connection, oh no!')

```

```

        let rawURL_arr = [reqData.rawURL]
        let startTimeArr = [startTime]
        clientProxyConn.on('data', (data) => {
            let tmpDetails =
getAddrAndPort(data.toString('binary'))
            rawURL_arr.push(tmpDetails.rawURL)
            startTimeArr.push(Date.now() /
1000)

            toServerConn.write(data)
        })
        // console.log(resData.toString())
        clientProxyConn.write(resData)
        let tmpRawURL = rawURL_arr.shift()
        if (timing) {
            // console.log({ url:
reqData.rawURL, cached: false, time: `${(Date.now() -
startTime).toString()}ms` })
            console.log({ url: tmpRawURL,
cached: false, time: `${(Date.now() - startTimeArr[0]).toString()}ms` })
            startTimeArr.shift()
        }
        addToCache(resData, tmpRawURL)

        //
clientProxyConn.pipe(toServerConn).pipe(clientProxyConn)

    }
    }
    })
} else {
    if (cachedRes.chunkArr) { //data was chunked, send
each TCP packet
        clientProxyConn.write(cachedRes.cachedStr)
        cachedRes.chunkArr.forEach(e => {
            clientProxyConn.write(e)
        })
    } else {
        // console.log(cachedRes.toString())
        clientProxyConn.write(cachedRes)
    }
    if (timing)
        console.log({ url: reqData.rawURL, cached:
true, time: `${(Date.now() - startTime).toString()}ms` })
    }
}

}

if (verbose)
    console.log({
        Message: 'Connection Established',
        Hostname: (reqData.isHTTPS ? reqData.host :
reqData.rawURL),
        Port: reqData.port,
        HTTPS: reqData.isHTTPS
    })

```

```

    })
    toServerConn.on('error', (err) => {
      if (!suppressErrs)
        console.error({ 'Server Error': err })
    })
    toServerConn.on('close', () => {
      console.warn({ 'Server Closed Conn':
`${reqData.host}:${reqData.port}` })
    })
  })
  clientProxyConn.on('error', (err) => {
    if (!suppressErrs)
      console.error({ 'Client Error': err })
  })
  clientProxyConn.on('close', () => {
    console.warn({ 'Client Closed Conn':
`${reqData.host}:${reqData.port}` })
  })
})
})

/**
 * Parses out: hostname, port and if a connection is HTTPS
 * @param {string} data data object stringified
 * @returns {{host:string, port:string, isHTTPS:boolean, rawURL:string}}
hostname, port, whether the connection is using HTTPS and the full path
trying to be accessed (if HTTP)
 */
let getAddrAndPort = (data) => {
  let hostData = []
  /*
   * Cannot actually read the data if using TLS but
   * HTTPS connections contain the keyword 'CONNECT'
   */
  hostData['isHTTPS'] = data.indexOf('CONNECT') !== -1
  if (hostData.isHTTPS) {
    let splitStr = data.split(` `)[1].split(`:`)
    hostData['host'] = splitStr[0]
    hostData['port'] = splitStr[1]
  } else {
    // console.log(data)
    hostData['rawURL'] = data.split('http://')[1].split(' ')[0]
    let splitStr = data.split(`Host: `)[1].split(`\r\n`)[0].split(`:`)
    hostData['host'] = splitStr[0]
    //HTTP defaults to port 80 but just in case...
    hostData['port'] = splitStr[1] ? splitStr[1] : '80'
  }
  return hostData
}

server.listen(port, () => {
  console.log(`Server running on: ${server.address().address !== '::' ?
server.address().address : 'localhost'}:${server.address().port}`)
})

```

```
stdin.addListener('data', (data) => {
  handleInput(data.toString().trim())
})

/**
 * Handles the strings input by the user
 *
 * @param {String} consoleInput The console input stringified and trimmed
 *
 * Commands:
 * block <domain> - adds domain to blocklist
 * unblock <domain> - removes domain from blocklist
 * cache - enables caching
 * nocache - disables caching
 * verbose - prints all connections to console
 * noverbose - disables printing of all connections to console
 * timing - print timing data of cache hits/misses
 * notiming - disables printing of cache hit/miss timing data
 * showsaving - shows how many bytes have been served from cache
 * showcachesize - shows how big the current cache is
 */
let handleInput = (consoleInput) => {
  let splitData = consoleInput.split(' ')
  let keyword = splitData[0]
  let param = splitData[1]
  switch (keyword) {
    case 'block':
      if (!blockList.blockedURLs.includes(param)) {
        blockList.blockedURLs.push(param)
        writeBlockList(blockList)
      } else {
        console.warn(`${param}, has already been blocked!`)
      }
      break
    case 'unblock':
      if (blockList.blockedURLs.includes(param)) {
        blockList.blockedURLs.splice(blockList.blockedURLs.indexOf(param), 1)
        writeBlockList(blockList)
      } else {
        console.warn(`${param}, was not blacklisted!`)
      }
      break
    case 'verbose':
      verbose = true
      break
    case 'noverbose':
      verbose = false
      break
    case 'cache':
      caching = true
      break
  }
}
```



```

        case 'nocache':
            caching = false
            break
        case 'timing':
            timing = true
            break
        case 'notiming':
            timing = false
            break
        case 'showsaving':
            console.log({ 'Saved Bytes':
bytesSavedFromNetwork.toLocaleString() })
            break
        case 'showcachesize':
            console.log({ 'Cache Size': currentCacheSize.toLocaleString()
})
            break
        default:
            console.error(`Input not recognised: ${keyword}, is not a
keyword`)
            break
    }
}

/**
 * Write updated blocklist
 * @param {blockList}
 */
let writeBlockList = (blockList) => {
    fs.writeFile(blockListPath, JSON.stringify(blockList), (err) => {
        if (err)
            console.error(`Could not write updated blocklist to disk!`)
        else
            console.log(`Updated blocklist written to disk`)
    })
}

/**
 * @param {string} url The requested URL
 * @return {Buffer|{cachedStr:Buffer, chunkArr:Array<Buffer>}}
 */
let getFromCache = (url) => {
    if (!caching)
        return false
    if (cache[url]) {
        let tmpCache = cache[url]
        if (tmpCache.expiryTime > (Date.now() / 1000)) {
            console.log(`Cached data for ${url}, found`)
            let cachedStr = tmpCache.firstHalfData +
(Math.floor(Date.now() / 1000) - tmpCache.startTime) +
tmpCache.secondHalfData

            cachedStr = Buffer.from(cachedStr, 'binary')
            // console.log({ cacheURL: url, cachedStr: cachedStr, str:

```

```

cachedStr.toString('binary') })
    bytesSavedFromNetwork += tmpCache.size
    //handle chunked and non-chunked data differently
    if (!tmpCache.chunkArr)
        return cachedStr
    return { cachedStr: cachedStr, chunkArr: tmpCache.chunkArr }
} else { //data is old
    console.log(`Cached data for ${url}, expired... purging`)
    cache = cache.splice(cache.indexOf('url'), 1)
    return false
}
}
return false
}

/**
 * @param {Buffer} responseBuffer The raw data response from the server
 * @param {string} url The url the request is for
 * @param {Array<Buffer>} chunkArr The chunks for a chunked response
 */
let addToCache = (responseBuffer, url, chunkArr = false) => {
    let parsedBuffer = responseBuffer.toString('binary')
    let parsedTotal = parsedBuffer.split('\r\n\r\n')
    let parsedBody = parsedTotal[1]
    let parsedBufferHead = parsedTotal[0] //extract the header
    if (parsedBufferHead.includes('404 Not Found\r\n')) { //dont cache 404
        // console.log(`Not caching 404 responses`)
        // return
    }
    if (parsedBufferHead.includes('Cache-Control: max-age=')) {
        let expiryTime = parsedBufferHead.split('Cache-Control: max-age=')[
1]

        if (expiryTime && expiryTime.split('\r\n', 1)[0]) {
            let size = responseBuffer.length
            if (chunkArr)
                chunkArr.forEach(e => {
                    size += e.length
                })
            expiryTime = expiryTime.split('\r\n', 1)[0].split(',')[0]
            expiryTime = parseInt(expiryTime) + Math.floor((Date.now() /
1000))

            if (parsedBufferHead.includes('Age: ')) { //header includes
age, this is ideal
                var ageSplit = parsedBufferHead.split('Age: ')

                var secondHalfData = ageSplit[1].split('\r\n')
                expiryTime -= parseInt(secondHalfData[0])
                var startTime = Math.floor(Date.now() / 1000) -
parseInt(secondHalfData[0])
                secondHalfData.splice(0, 1)
                secondHalfData = secondHalfData.join('\r\n') + '\r\n\r\n'
+ parsedBody.toString('binary')
            } else { //header does not include age, this is not ideal
                let theHeaderArr = parsedBufferHead.split('Cache-Control:

```

```

max-age=')
    var startTime = 0
    theHeaderArr[0] += 'Cache-Control: max-age='
    theHeaderArr[1] = theHeaderArr[1].split('\r\n')
    theHeaderArr[0] += theHeaderArr[1][0] + '\r\n'
    var ageSplit = []
    ageSplit[0] = theHeaderArr[0]
    var secondHalfData =
parsedBuffer.slice(ageSplit[0].length)
    // console.log({ secondHalf: secondHalfData })
}

    cache[url] = {
        expiryTime: expiryTime,
        firstHalfData: ageSplit[0] + 'Age: ',
        secondHalfData: '\r\n' + secondHalfData,
        startTime: startTime,
        chunkArr: chunkArr,
        size: size
    }
    currentCacheSize += size
    console.log({ CachedURL: url, Size: `${size.toLocaleString()}
bytes` })
    } else {
        // console.log(parsedBufferHead)
        console.log(`Could not cache response from: ${url}, due to
header parameters`)
    }
    } else {
        // console.log(parsedBufferHead)
        console.log(`Could not cache response from: ${url}, due to header
parameters`)
    }
}

/**
 * Determines if a HTTP request is for a websocket
 * @param {Buffer} rawData The raw request data
 * @param {Boolean} allowNoCache Respond true to no cache requests
 */
let isWebSocketRequest = (rawData) => {
    let stringifiedData = rawData.toString()
    if (stringifiedData.includes('Upgrade: websocket\r\n') ||
stringifiedData.includes('Connection: upgrade\r\n'))
        return true
    return false
}

/**
 * Determines if a response wants to be cached
 * @param {Buffer} rawData
 */
let isCacheableResponse = (rawData) => {
    let stringifiedData = rawData.toString()

```

```
    if (stringifiedData.includes('Cache-Control: no-cache\r\n') ||  
        stringifiedData.includes('Pragma: no-cache\r\n') ||  
        !stringifiedData.includes('Cache-Control: max-age='))  
        return false  
    return true  
}
```